

Introduction to Databases & SQL

T5 Bootcamp by SDAIA



SDAIA
الهيئة السعودية للبيانات
والذكاء الاصطناعي
Saudi Data & AI Authority

Introduction to Databases & SQL



SDAIA
الهيئة السعودية للبيانات
والذكاء الاصطناعي
Saudi Data & AI Authority

Agenda

Program layout

- Types of Databases and Database Applications
- Basic Definitions
- Typical DBMS Functionality
- Example of a Database (UNIVERSITY/Company)
- Main Characteristics of the Database Approach
- Advantages of Using the Database Approach
- When Not to Use Databases
- Structured Query Language



Environment

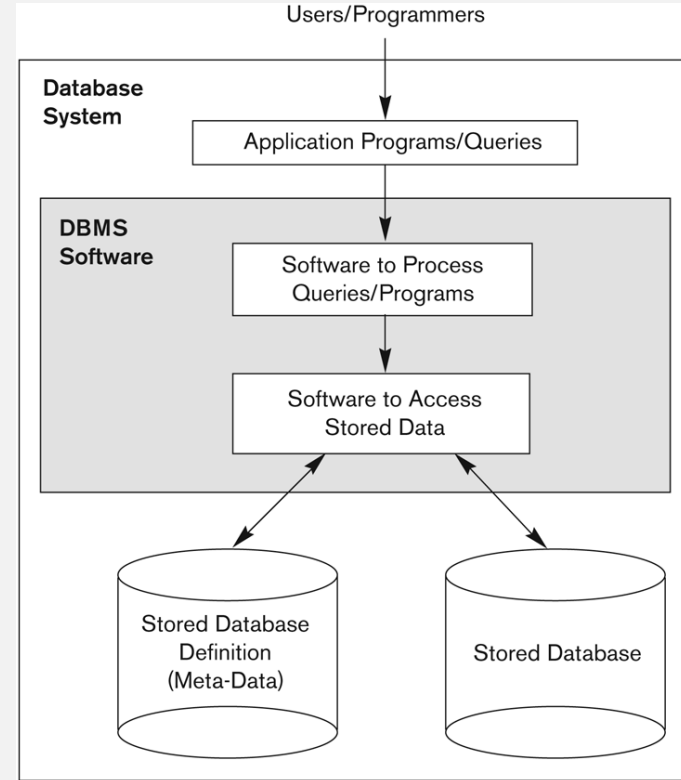
- Grading system
- Module Architecture
 - Slides
 - Hands-on Practice
 - Project
 - Task
 - Exam
- Oracle Live SQL





Basic Definitions


- **Database:** collection of related data.
- **Data:** Known facts that can be recorded and have an implicit meaning.
- **Database Management System (DBMS):** A software package/ system to facilitate the creation and maintenance of a computerized database.
- **Database System:** The DBMS software together with the data itself. Sometimes, the applications are also included.





Types of Databases and Database Applications

- Traditional Applications:
 - Numeric and Textual Databases
- Other Databases:
 - Multimedia Databases
 - Geographic Information Systems (GIS)
 - Data Warehouses
 - Real-time and Active Databases
 - Many other applications
 - NoSql Databases





DBMS Functionality

- Specify the characteristics of a specific database, including its data types, structures, and constraints.
- Populate or Create the initial contents of the database on a secondary storage device.
- Interacting with the database:
 - Retrieval: Performing queries and generating reports.
 - Modification: Making insertions, deletions, and updates to its data.
 - Accessing the database via web applications.
- Managing and distributing data among multiple users and application programs concurrently while ensuring data integrity and consistency.





Example of University Database

- Mini-world: University
- Some mini-world entities:
 - STUDENTs
 - COURSEs
 - SECTIONs (of COURSEs)
 - (academic) DEPARTMENTs
 - INSTRUCTORs
- Relationships:
 - SECTIONs are of specific COURSEs
 - STUDENTs take SECTIONs
 - COURSEs have prerequisite COURSEs
 - INSTRUCTORs teach SECTIONs
 - COURSEs are offered by DEPARTMENTs
 - STUDENTs major in DEPARTMENTs

COURSE

| Course_name | Course_number | Credit_hours | Department |
|---------------------------|---------------|--------------|------------|
| Intro to Computer Science | CS1310 | 4 | CS |
| Data Structures | CS3320 | 4 | CS |
| Discrete Mathematics | MATH2410 | 3 | MATH |
| Database | CS3380 | 3 | CS |

SECTION

| Section_identifier | Course_number | Semester | Year | Instructor |
|--------------------|---------------|----------|------|------------|
| 85 | MATH2410 | Fall | 04 | King |
| 92 | CS1310 | Fall | 04 | Anderson |
| 102 | CS3320 | Spring | 05 | Knuth |
| 112 | MATH2410 | Fall | 05 | Chang |
| 119 | CS1310 | Fall | 05 | Anderson |
| 135 | CS3380 | Fall | 05 | Stone |

GRADE_REPORT

| Student_number | Section_identifier | Grade |
|----------------|--------------------|-------|
| 17 | 112 | B |
| 17 | 119 | C |
| 8 | 85 | A |
| 8 | 92 | A |
| 8 | 102 | B |
| 8 | 135 | A |

PREREQUISITE

| Course_number | Prerequisite_number |
|---------------|---------------------|
| CS3380 | CS3320 |
| CS3380 | MATH2410 |
| CS3320 | CS1310 |



Main Characteristics of the Database Approach

- The self-descriptive aspect of a database system: DBMS catalog and metadata.
- Separation between programs and data (program-data independence).
- Data Abstraction: Utilization of a data model.
- Provision for multiple perspectives of the data.
- Facilitation of data sharing and multi-user transaction processing (support for concurrent users).

Advantages of Using the Database Approach

- Managing data redundancy effectively.
- Enabling data sharing among multiple users.
- Implementing measures to prevent unauthorized access to data.
- Ensuring persistent storage for object-oriented program objects.
- Establishing storage structures to optimize query processing efficiency.
- Offering backup and recovery functionalities.
- Offering various interfaces tailored to different user groups.
- Modeling intricate relationships within the data.
- Enforcing integrity constraints on the database.
- Deriving insights and executing actions based on stored data using deductive and active rules.





When not to use a DBMS?

- Main challenges (costs) associated with using a DBMS:
 - High initial investment and potential need for additional hardware.
 - Overhead related to providing generality, security, concurrency control, recovery, and integrity functions.
- Instances where a DBMS might be unnecessary:
 - In cases where the database and applications are straightforward, well-defined, and unlikely to change.
 - When stringent real-time requirements cannot be met due to DBMS overhead.
 - If there is no requirement for multiple users to access the data concurrently.
- Instances where no DBMS may suffice:
 - When the complexity of the data exceeds the capabilities of the database system due to modeling limitations.
 - If database users require special operations that are not supported by the DBMS.



Data Modeling



SDAIA
الهيئة السعودية للبيانات
والذكاء الاصطناعي
Saudi Data & AI Authority



Data Models

- Data Model:
 - A set of concepts to describe database structure, the operations for manipulating these structures, and certain database enforced constraints.
- Data Model Structure and Constraints:
 - Constructs are used to define the database structure
 - Constructs typically include elements (and their data types) as well as groups of elements (e.g. entity, record, table), and relationships among such groups
 - Constraints specify some restrictions on valid data; these constraints must be enforced at all times



Schemas versus Instances

- Database Schema:
 - The **description** of a database structure, data types, and the constraints on the database.
- Schema Diagram:
 - Illustrative** display of a database schema.
- Schema Construct:
 - A **component** of the schema or an object within the schema, e.g., STUDENT, COURSE.

| STUDENT | | | |
|---------|----------------|-------|-------|
| Name | Student_number | Class | Major |

Schema database

| COURSE | | | |
|-------------|---------------|--------------|------------|
| Course_name | Course_number | Credit_hours | Department |

| PREREQUISITE | |
|---------------|---------------------|
| Course_number | Prerequisite_number |

| SECTION | | | | |
|--------------------|---------------|----------|------|------------|
| Section_identifier | Course_number | Semester | Year | Instructor |

| GRADE_REPORT | | |
|----------------|--------------------|-------|
| Student_number | Section_identifier | Grade |

| COURSE | | | |
|---------------------------|---------------|--------------|------------|
| Course_name | Course_number | Credit_hours | Department |
| Intro to Computer Science | CS1310 | 4 | CS |
| Data Structures | CS3320 | 4 | CS |
| Discrete Mathematics | MATH2410 | 3 | MATH |
| Database | CS3380 | 3 | CS |

| SECTION | | | | |
|--------------------|---------------|----------|------|------------|
| Section_identifier | Course_number | Semester | Year | Instructor |
| 85 | MATH2410 | Fall | 04 | King |
| 92 | CS1310 | Fall | 04 | Anderson |
| 102 | CS3320 | Spring | 05 | Knuth |
| 112 | MATH2410 | Fall | 05 | Chang |
| 119 | CS1310 | Fall | 05 | Anderson |
| 135 | CS3380 | Fall | 05 | Stone |

| GRADE_REPORT | | |
|----------------|--------------------|-------|
| Student_number | Section_identifier | Grade |
| 17 | 112 | B |
| 17 | 119 | C |
| 8 | 85 | A |
| 8 | 92 | A |
| 8 | 102 | B |
| 8 | 135 | A |

| PREREQUISITE | |
|---------------|---------------------|
| Course_number | Prerequisite_number |
| CS3380 | CS3320 |
| CS3380 | MATH2410 |
| CS3320 | CS1310 |





DBMS Languages - SQL

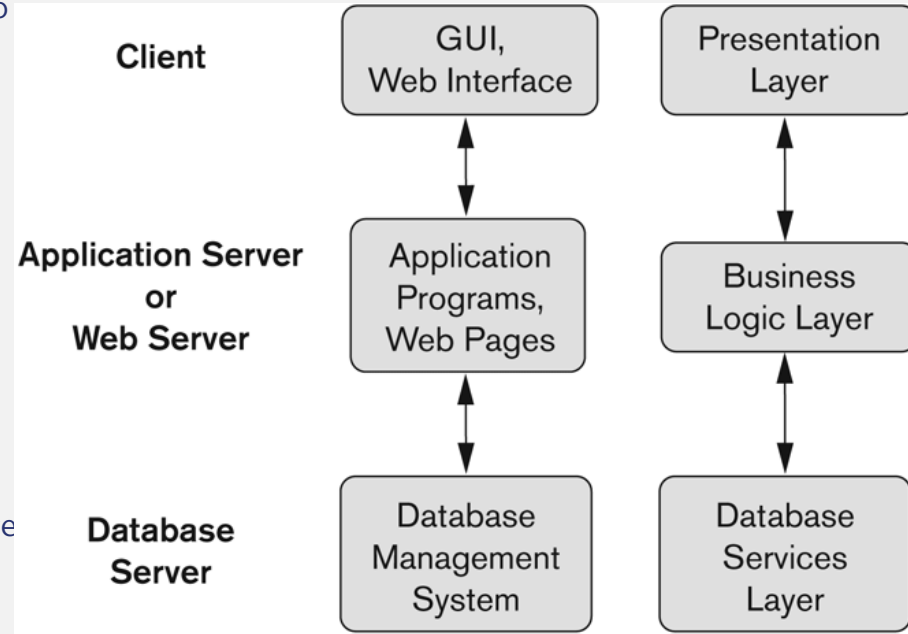
- **Data Definition Language (DDL)**
 - Used for defining and managing the structure of database objects. It includes commands for creating, altering, and dropping database objects such as tables, indexes, views, and constraints.
 - Examples of DDL commands include CREATE, ALTER, DROP, and TRUNCATE.
- **Data Manipulation Language (DML)**
 - Used for managing data within a database. It includes commands for inserting, updating, deleting, and retrieving data from database tables.
 - Examples of DML commands include INSERT, UPDATE, DELETE, and SELECT.
- **Data Control Language (DCL)**
 - Used for controlling access to data within a database. It includes commands for granting and revoking permissions and privileges to database objects.
 - Examples of DCL commands include GRANT and REVOKE.





DBMS Server

- Furnishes database inquiry and transaction utilities to the users.
- Relational Database Management System (RDBMS) servers are commonly referred to as SQL servers, query servers, or transaction servers.
- Client applications leverage an Application Programming Interface (API) to connect to server databases through a standard interface, such as:
 - ODBC: Open Database Connectivity standard
 - JDBC: Java Database Connectivity for Java programming access.
- Both the client and server necessitate installation of the suitable client module and server module software for ODBC or JDBC.





Classification of DBMSs

- Based on the data model used
 - Traditional: Relational, Network, Hierarchical.
 - Emerging: Object-oriented, Object-relational.
- Other classifications
 - Single-user (personal computers) vs. multi-user (most DBMSs).
 - Centralized (uses a single computer with one database) vs. distributed (uses multiple computers, multiple databases)

Variations of Distributed DBMSs (DDBMSs)

- Homogeneous DDBMS
- Heterogeneous DDBMS
- Federated or Multidatabase Systems



Data Modeling Using the Entity-Relationship (ER) Model



SDAIA
الهيئة السعودية للبيانات
والذكاء الاصطناعي
Saudi Data & AI Authority

Agenda

- Overview of Database Design Process
- Example Database Application (COMPANY)
- ER Model Concepts
 - Entities and Attributes
 - Entity Types, Value Sets, and Key Attributes
 - Relationships and Relationship Types
 - Weak Entity Types
- Roles and Attributes in Relationship Types
- ER Diagrams - Notation
- ER Diagram for COMPANY Schema





Use Case: COMPANY Database

We need to create a database schema design based on the following requirements of the COMPANY Database:

- The company is organized into DEPARTMENTS. Each department has a name, number and an employee who manages the department. We keep track of the start date of the department manager. A department may have several locations.
- Each department controls a number of PROJECTs. Each project has a unique name, unique number and is located at a single location.
- We store each EMPLOYEE's social security number, address, salary, sex, and birthdate.
 - Each employee works for one department but may work on several projects.
 - We keep track of the number of hours per week that an employee currently works on each project.
 - We also keep track of the direct supervisor of each employee.
- Each employee may have a number of DEPENDENTs.
 - For each dependent, we keep track of their name, sex, birthdate, and relationship to the employee.





Entities & Attributes


- Entities refer to distinct objects or entities within the real-world domain, which are represented within the database.
 - For instance, an EMPLOYEE named John Smith, the Research DEPARTMENT, or the ProductX PROJECT.
- Attributes serve as characteristics utilized to describe an entity.
 - For instance, an EMPLOYEE entity may possess attributes such as Name, SSN, Address, Sex, and BirthDate.
- Each entity possesses a unique set of values corresponding to its attributes.
 - For instance, a specific employee entity may have attributes with values such as Name='John Smith', SSN='123456789', Address='731 Fondren, Houston, TX', Sex='M', and BirthDate='09-JAN-55'.
- Every attribute is associated with a specific value set or data type, such as integer, string, subrange, enumerated type, etc.





Types of Attributes

- Simple
 - Single atomic value for the attribute. SSN or Sex.
- Composite
 - Attribute may be composed of several components.
For example:
 - Address(Apt#, House#, Street, City, State, ZipCode, Country), or
 - Name(FirstName, MiddleName, LastName).
- Multi-valued
 - An attribute may have multiple values. For example, Color of a CAR or PreviousDegrees of a STUDENT.





Entity Types and Key Attributes

- Entities with the same basic attributes are grouped or typed into an **entity type**. Ex: entity type EMPLOYEE and PROJECT.
- An attribute of an entity type for which each entity must have a unique value is called a **key attribute** of the entity type. Ex: SSN of EMPLOYEE.






Entity Types and Key Attributes

- A key attribute may be composite.
 - VehicleTagNumber is a key of the CAR entity type with components (Number, State).
- An entity type may have more than one key.
 - The CAR entity type may have two keys:
VehicleIdentificationNumber (popularly called VIN)
VehicleTagNumber (Number, State), aka license plate number.
- Each key is underlined



Displaying an Entity in ER Diagram

- Entity type is displayed in a rectangular box
- Attributes are displayed in ovals
 - Each attribute is connected to its entity type
 - Components of a composite attribute are connected to the oval representing the composite attribute
 - Each key attribute is underlined
 - Multivalued attributes displayed in double ovals



CAR
Registration (Number, State), Vehicle_id, Make, Model, Year, {Color}

CAR₁
((ABC 123, TEXAS), TK629, Ford Mustang, convertible, 2004 {red, black})

CAR₂
((ABC 123, NEW YORK), WP9872, Nissan Maxima, 4-door, 2005, {blue})

CAR₃
((VSY 720, TEXAS), TD729, Chrysler LeBaron, 4-door, 2002, {white, blue})

⋮



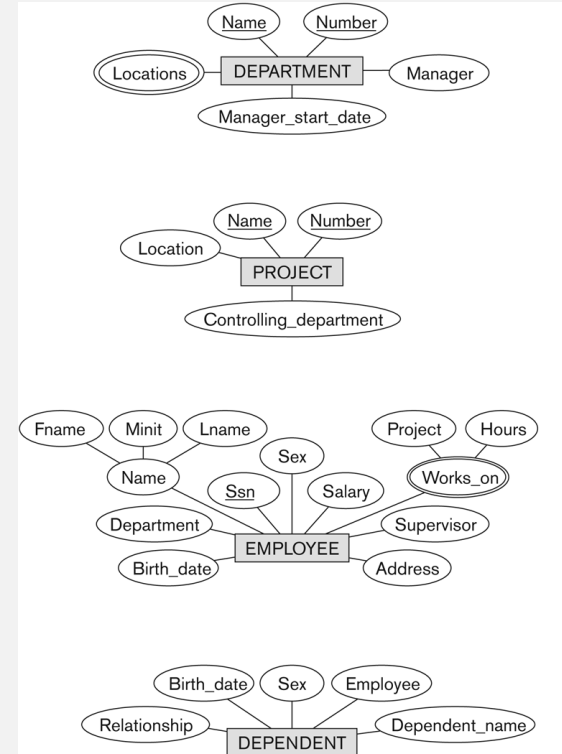


Initial Design of Entity Types: EMPLOYEE, DEPARTMENT, PROJECT, DEPENDENT

Based on the requirements, we can identify four initial entity types in the COMPANY database:

- DEPARTMENT
- PROJECT
- EMPLOYEE
- DEPENDENT

Their initial design is shown on the following slide
The initial attributes shown are derived from the requirements description





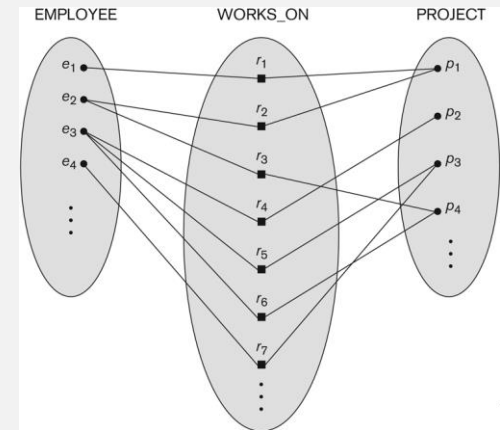
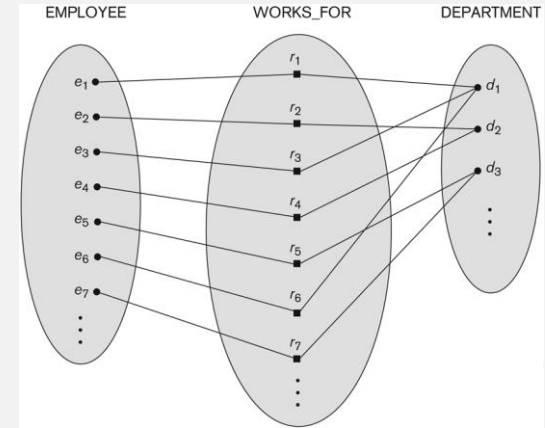
Introducing relationships

ER model three main concepts: Entities, Attributes, Relationships
Relationship relates two or more distinct entities with a specific meaning.

- EMPLOYEE John Smith works on the ProductX PROJECT
- EMPLOYEE Franklin Wong manages the Research DEPARTMENT.

Relationship degree is the number of participating entity types.

- Ex: MANAGES and WORKS_ON are binary relationships.

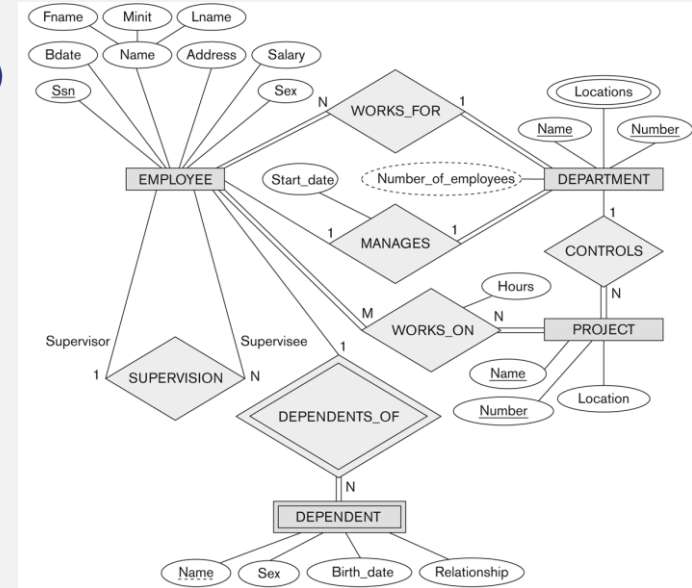




Refining the COMPANY database schema by introducing relationships

Six relationship types are identified

- WORKS_FOR (between EMPLOYEE, DEPARTMENT)
- MANAGES (between EMPLOYEE, DEPARTMENT)
- CONTROLS (between DEPARTMENT, PROJECT)
- WORKS_ON (between EMPLOYEE, PROJECT)
- SUPERVISION (between EMPLOYEE (as subordinate), EMPLOYEE (as supervisor))
- DEPENDENTS_OF (between EMPLOYEE, DEPENDENT)





Recursive Relationship Type (Self Join)


Relationship with the same participating entity type in distinct roles

Example: the SUPERVISION relationship

EMPLOYEE participates twice in two distinct roles:

- supervisor (or boss) role
- supervisee (or subordinate) role

Can you give another example?





Weak Entity Types

Entity without a key attribute and participate in an identifying relationship type with an owner or identifying entity

Entities are identified by the combination of:

- A partial key of the weak entity type
- The particular entity they are related to in the identifying entity type

Example:

- A DEPENDENT entity is identified by the dependent's first name, and the specific EMPLOYEE with whom the dependent is related
- Name of DEPENDENT is the partial key
- DEPENDENT is a weak entity type
- EMPLOYEE is its identifying entity type via the identifying relationship type DEPENDENT_OF





Relationships Constraints

Cardinality Ratio (specifies maximum participation)

- One-to-one (1:1)
- One-to-many (1:N) or Many-to-one (N:1)
- Many-to-many (M:N)

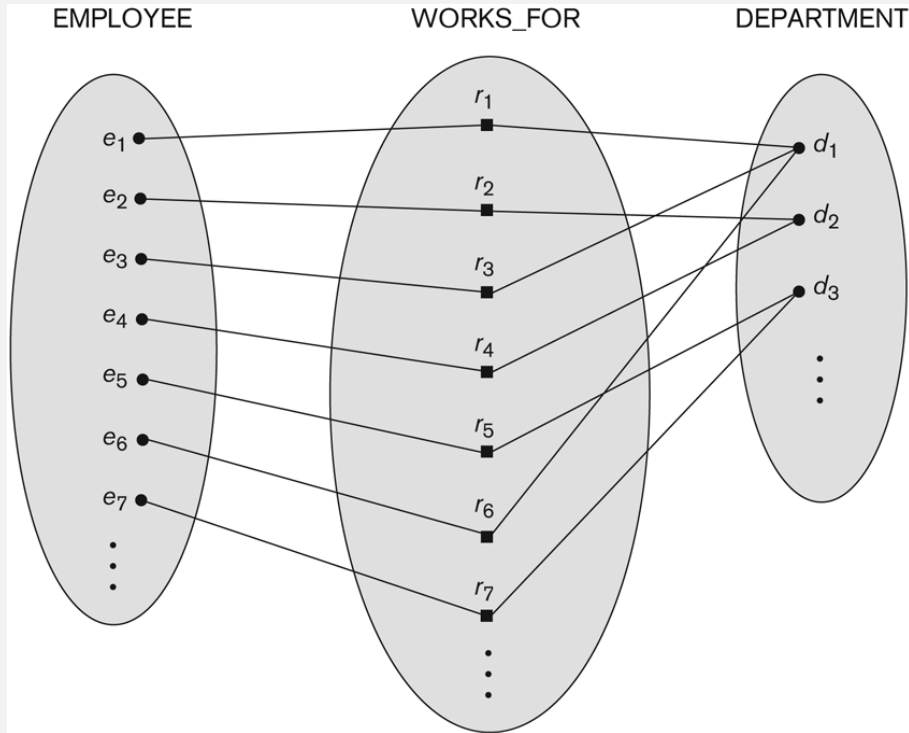
Existence Dependency Constraint (specifies minimum participation)

- zero (optional participation, not existence-dependent)
- one or more (mandatory participation, existence-dependent)

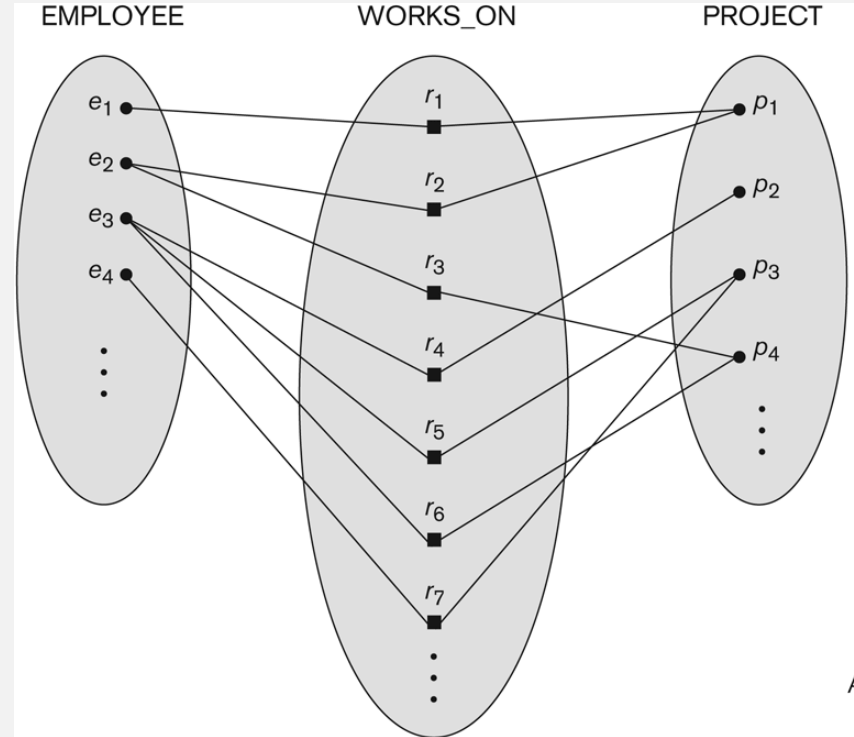




Many-to-one (N:1) Relationship



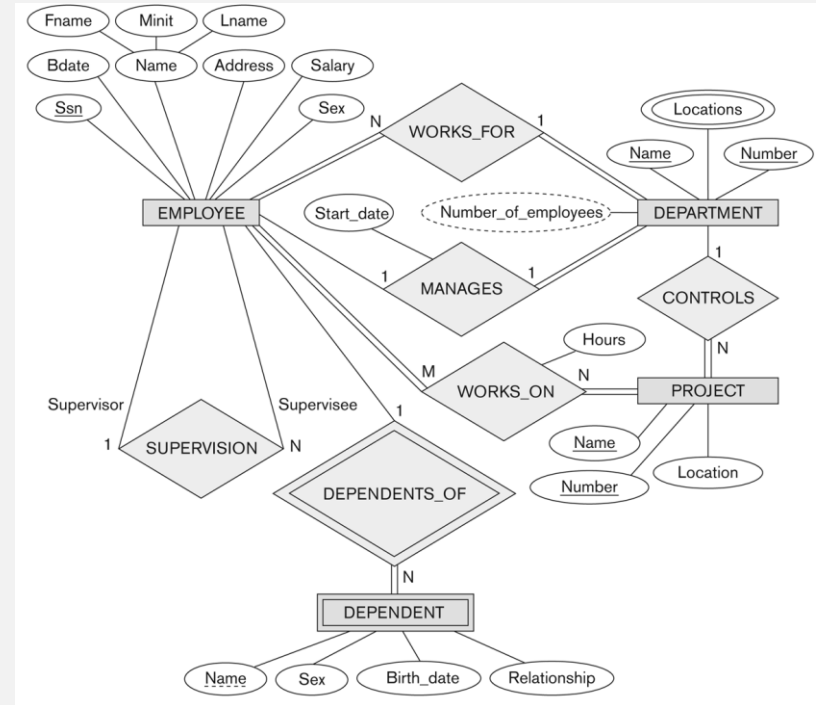
Many-to-many (M:N) Relationship




Relationship Attributes

A relationship type can have attributes:

- Ex: HoursPerWeek of WORKS_ON
- Its value for each relationship instance describes the number of hours per week that an EMPLOYEE works on a PROJECT.
- Mostly used with M:N relationships

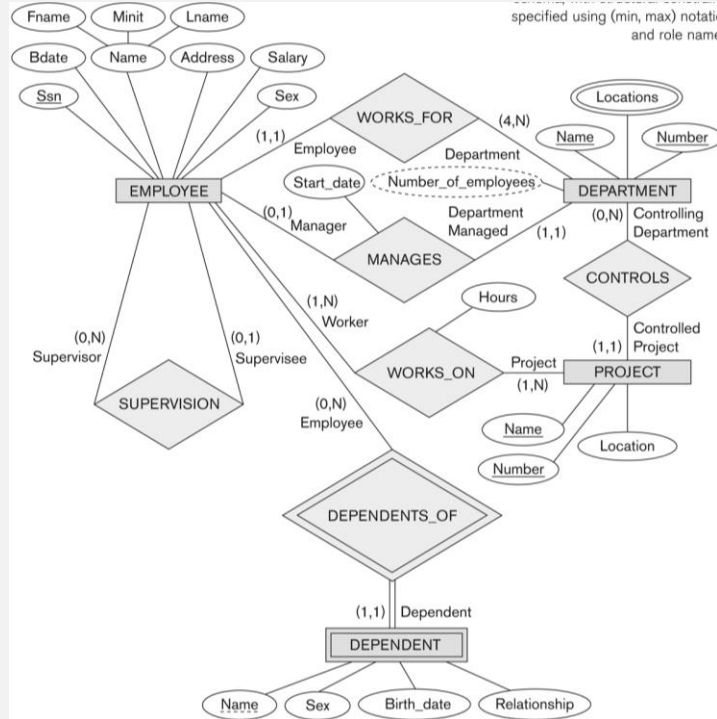


(min,max) notation





COMPANY ER Schema Diagram with (min, max) notation



| Symbol | Meaning |
|--------|---|
| | Entity |
| | Weak Entity |
| | Relationship |
| | Identifying Relationship |
| | Attribute |
| | Key Attribute |
| | Multivalued Attribute |
| | Composite Attribute |
| | Derived Attribute |
| | Total Participation of E_2 in R |
| | Cardinality Ratio 1: N for $E_1:E_2$ in R |
| | Structural Constraint (min, max) on Participation of E in R |



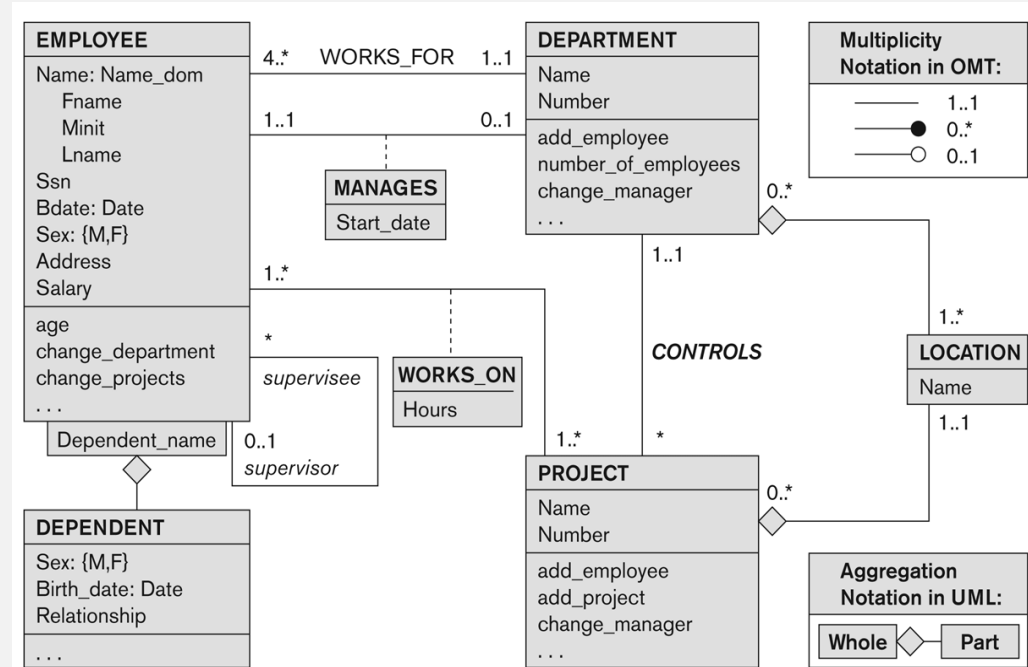


UML class diagrams

Represent classes (entity types) as large rounded boxes

Three sections:

- Entity type (class) name
- Attributes,
- Class operations





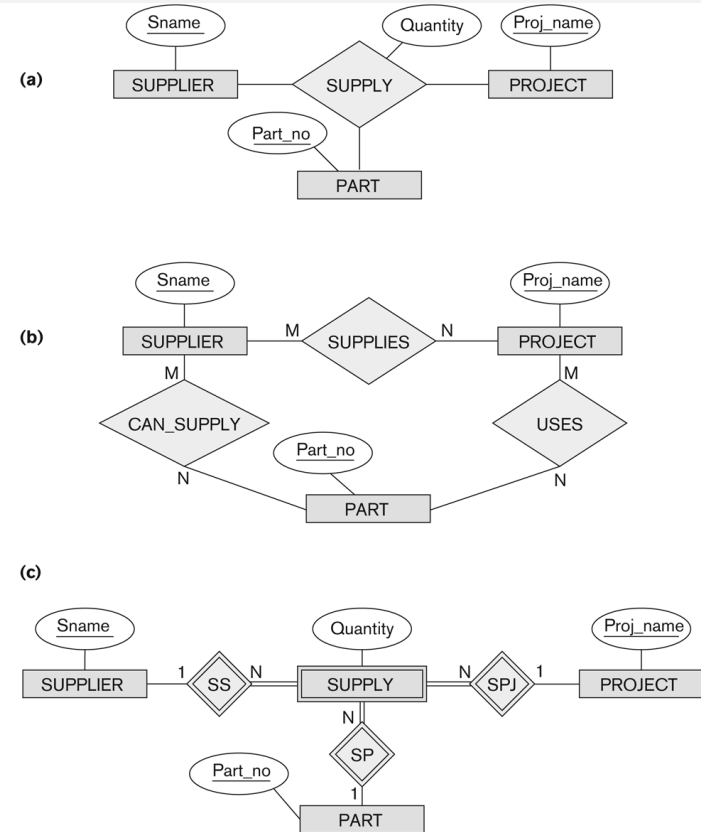
Higher Degree Relationships

Binary Relationship: degree 2

Ternary Relationship: degree 3

N-ary: degree N

Degree specifies number of entities participating in a relationship



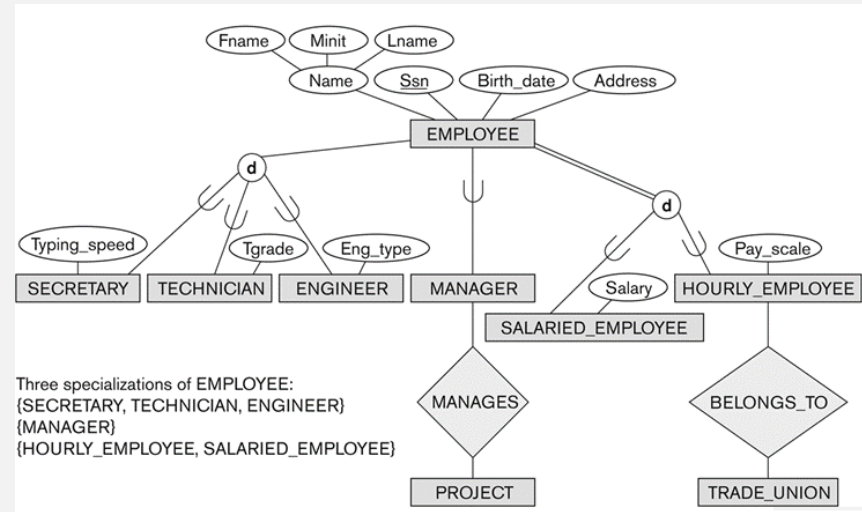


Subclasses and Superclasses

Subgroupings entities

Disjointness Constraint:

- Disjoint (d): an entity can be a member of at most one of the subclasses
- Overlapping (o): that is the same entity may be a member of more than one subclasses





Relational Model Concepts

- The relational Model of Data is based on the concept of a Relation
- The model was first proposed by Dr. E.F. Codd of IBM Research in 1970 in ACM paper





Informal Definitions

Relation is a table of values.

Relation contains a set of rows.

Data elements in each row corresponds to a real-world entity or relationship (tuples)

Column indicates data items meaning in that column (attribute)

The diagram illustrates the components of a relation. The label 'Relation Name' points to the word 'STUDENT' above the table. The label 'Attributes' points to the column headers: 'Name', 'Ssn', 'Home_phone', 'Address', 'Office_phone', 'Age', and 'Gpa'. The label 'Tuples' points to the rows of data in the table.

| Name | Ssn | Home_phone | Address | Office_phone | Age | Gpa |
|----------------|-------------|------------|----------------------|--------------|-----|------|
| Benjamin Bayer | 305-61-2435 | 373-1616 | 2918 Bluebonnet Lane | NULL | 19 | 3.21 |
| Chung-cha Kim | 381-62-1245 | 375-4409 | 125 Kirby Road | NULL | 18 | 2.89 |
| Dick Davidson | 422-11-2320 | NULL | 3452 Elgin Road | 749-1253 | 25 | 3.53 |
| Rohan Panchal | 489-22-1100 | 376-9821 | 265 Lark Lane | 749-6492 | 28 | 3.93 |
| Barbara Benson | 533-69-1238 | 839-8461 | 7384 Fontana Lane | NULL | 19 | 3.25 |





Relational Integrity Constraints

Constraints are conditions that must hold on all valid relation states.

There are three main types of constraints in the relational model:

- **Key** constraints
- **Entity integrity** constraints
- **Referential integrity** constraints
- **Domain** constraint

Every value in a tuple must be from the domain of its attribute (or it could be null, if allowed for that attribute)

- Several candidate keys, choose one to be your **primary key**.
- Primary key attributes (**underlined**)

CAR

| <u>License_number</u> | Engine_serial_number | Make | Model | Year |
|-----------------------|----------------------|------------|---------|------|
| Texas ABC-739 | A69352 | Ford | Mustang | 02 |
| Florida TVP-347 | B43696 | Oldsmobile | Cutlass | 05 |
| New York MPO-22 | X83554 | Oldsmobile | Delta | 01 |
| California 432-TFY | C43742 | Mercedes | 190-D | 99 |
| California RSK-629 | Y82935 | Toyota | Camry | 04 |
| Texas RSK-629 | U028365 | Jaguar | XJS | 04 |





COMPANY Database Schema

EMPLOYEE

| | | | | | | | | | |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|
| Frame | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| | | | |
|-------|----------------|---------|----------------|
| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|-------|----------------|---------|----------------|

DEPT_LOCATIONS

| | |
|----------------|------------------|
| <u>Dnumber</u> | <u>Dlocation</u> |
|----------------|------------------|

PROJECT

| | | | |
|-------|----------------|-----------|------|
| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-------|----------------|-----------|------|

WORKS_ON

| | | |
|-------------|------------|-------|
| <u>Essn</u> | <u>Pno</u> | Hours |
|-------------|------------|-------|

DEPENDENT

| | | | | |
|-------------|-----------------------|-----|-------|--------------|
| <u>Essn</u> | <u>Dependent_name</u> | Sex | Bdate | Relationship |
|-------------|-----------------------|-----|-------|--------------|



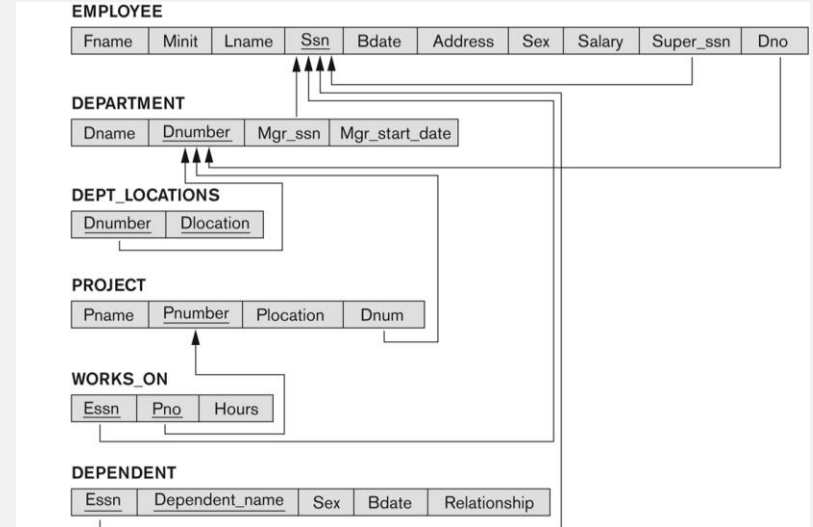
Entity Integrity vs. Referential Integrity

Entity Integrity:

- Primary key cannot have null values

Referential Integrity:

- Specifies a relationship among tuples in two relations



Populated database state for COMPANY

EMPLOYEE

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|-----|-------|---------|-----|--------|-----------|-----|
|-------|-------|-------|-----|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|-------|---------|---------|----------------|
|-------|---------|---------|----------------|

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
|---------|-----------|

PROJECT

| Pname | Pnumber | Plocation | Dnum |
|-------|---------|-----------|------|
|-------|---------|-----------|------|

WORKS_ON

| Essn | Pno | Hours |
|------|-----|-------|
|------|-----|-------|

DEPENDENT

| Essn | Dependent_name | Sex | Bdate | Relationship |
|------|----------------|-----|-------|--------------|
|------|----------------|-----|-------|--------------|

EMPLOYEE

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------------|---------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

WORKS_ON

| Essn | Pno | Hours |
|-----------|-----|-------|
| 123456789 | 1 | 32.5 |
| 123456789 | 2 | 7.5 |
| 666884444 | 3 | 40.0 |
| 453453453 | 1 | 20.0 |
| 453453453 | 2 | 20.0 |
| 333445555 | 2 | 10.0 |
| 333445555 | 3 | 10.0 |
| 333445555 | 10 | 10.0 |
| 333445555 | 20 | 10.0 |
| 999887777 | 30 | 30.0 |
| 999887777 | 10 | 10.0 |
| 987987987 | 10 | 35.0 |
| 987987987 | 30 | 5.0 |
| 987654321 | 30 | 20.0 |
| 987654321 | 20 | 15.0 |
| 888665555 | 20 | NULL |

PROJECT

| Pname | Pnumber | Plocation | Dnum |
|-----------------|---------|-----------|------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Stafford | 4 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

DEPENDENT

| Essn | Dependent_name | Sex | Bdate | Relationship |
|-----------|----------------|-----|------------|--------------|
| 333445555 | Alice | F | 1986-04-05 | Daughter |
| 333445555 | Theodore | M | 1983-10-25 | Son |
| 333445555 | Joy | F | 1958-05-03 | Spouse |
| 987654321 | Abner | M | 1942-02-28 | Spouse |
| 123456789 | Michael | M | 1988-01-04 | Son |
| 123456789 | Alice | F | 1988-12-30 | Daughter |
| 123456789 | Elizabeth | F | 1967-05-05 | Spouse |



Update Operations on Relations (Tables)

INSERT a tuple (row).

DELETE a tuple (row).

MODIFY a tuple (row).

Meanwhile:

- Integrity constraints are not violated.
- You can group several update operations together.
- Updates are propagated automatically (integrity constraints).

- Integrity violation:
 - Operation is cancelled (RESTRICT or REJECT option)
 - Inform the user and perform the operation.
 - Additional updates are triggered to correct violations (CASCADE option, SET NULL option)
 - Execute a user-specified error-correction routine





Possible violations:

- INSERT new tuple:
 - Domain constraint:
 - New tuple value is not of the specified attribute domain
 - Key constraint:
 - Key value in the new tuple already
 - Referential integrity:
 - Foreign key value does not exist in the referenced relation
 - Entity integrity:
 - Primary key value is null
- DELETE :
 - Primary key value is referenced from other tuples in the database
 - Can be remedied by several actions: RESTRICT, CASCADE, SET NULL RESTRICT option: reject the deletion
 - Above options must be specified during foreign key constraint design





Possible violations

- UPDATE:
Domain constraint and NOT NULL constraint
Other constraints may also be violated:
 - Updating primary key (PK):
 - Similar to a DELETE followed by an INSERT
 - Specify similar options to DELETE
 - Updating foreign key (FK):
 - May violate referential integrity
 - Updating an ordinary attribute:
 - Can only violate domain constraints





In-Class Exercise

Consider the following relations for a database that keeps track of student enrollment in courses and the books adopted for each course:

STUDENT(SSN, Name, Major, Bdate)

COURSE(Course#, Cname, Dept)

ENROLL(SSN, Course#, Quarter, Grade)

BOOK_ADOPTION(Course#, Quarter, Book_ISBN)

TEXT(Book_ISBN, Book_Title, Publisher, Author)

Draw a relational schema diagram specifying the foreign keys for this schema.



SQL

Structured Query Language



SDAIA

الهيئة السعودية للبيانات
والذكاء الاصطناعي
Saudi Data & AI Authority



CREATE TABLE

- Creates a new base table with a name, attributes and their data types (INTEGER, FLOAT, DECIMAL(i,j), CHAR(n), VARCHAR(n))

Specify NOT NULL constraint

```
CREATE TABLE DEPARTMENT (  
  DNAME    VARCHAR(10) NOT NULL,  
  DNUMBER  INTEGER      NOT NULL,  
  MGRSSN   CHAR(9),  
  MGRSTARTDATE CHAR(9)  
  );
```

- Primary key attributes, secondary keys, and referential integrity constraints

```
CREATE TABLE DEPT (  
  DNAME VARCHAR(10)    NOT NULL,  
  DNUMBER INTEGER      NOT NULL,  
  MGRSSN CHAR(9),  
  MGRSTARTDATE CHAR(9),  
  PRIMARY KEY (DNUMBER),  
  UNIQUE (DNAME),  
  FOREIGN KEY (MGRSSN) REFERENCES  
  EMP  
  );
```





DROP TABLE

DROP TABLE DEPENDENT ;

ALTER TABLE

- Adds an attribute:
 - NULLs in all the previous existing (NOT NULL constraint not allowed)

**ALTER TABLE EMPLOYEE
ADD JOB VARCHAR(12) ;**





REFERENTIAL INTEGRITY OPTIONS

- Specify RESTRICT, CASCADE, SET NULL or SET DEFAULT

```
CREATE TABLE DEPT (  
    DNAME          VARCHAR(10)          NOT NULL,  
    DNUMBER        INTEGER              NOT NULL,  
    MGRSSN         CHAR(9) ,  
    MGRSTARTDATE   CHAR(9) ,  
    PRIMARY KEY (DNUMBER) ,  
    UNIQUE (DNAME) ,  
    FOREIGN KEY (MGRSSN) REFERENCES EMP  
    ON DELETE SET DEFAULT ON UPDATE CASCADE) ;
```





REFERENTIAL INTEGRITY OPTIONS

```
CREATE TABLE EMP (  
    ENAME          VARCHAR(30)  NOT NULL,  
    ESSN           CHAR(9) ,  
    BDATE          DATE ,  
    DNO            INTEGER  DEFAULT 1 ,  
    SUPERSSN       CHAR(9) ,  
    PRIMARY KEY (ESSN) ,  
    FOREIGN KEY (DNO) REFERENCES DEPT  
    ON DELETE SET DEFAULT ON UPDATE CASCADE ,  
    FOREIGN KEY (SUPERSSN) REFERENCES EMP ON DELETE  
    SET NULL ON UPDATE CASCADE) ;
```





Additional Data Types in SQL2 and SQL-99

DATE, TIME, and TIMESTAMP data types

DATE:

- format yyyy-mm-dd

TIME:

- format hh:mm:ss

TIME(i):

- Specifies fractions of a second
- format is hh:mm:ss.ii...i

● TIMESTAMP:

- DATE and TIME components

● INTERVAL:

- Relative value (not absolute value)
- DAY/TIME intervals
- YEAR/MONTH intervals
- Positive or negative
- Added to or subtracted from an absolute value, the result is an absolute value





Retrieval Queries in SQL

Basic statement : SELECT statement

Basic form of the SQL SELECT statement is called a mapping or a SELECT-FROM-WHERE block

| | | |
|---------------|--------------------|-------------------------------|
| SELECT | [DISTINCT] | <attribute list> |
| FROM | | <table list> |
| WHERE | | <condition> |

- <attribute list> : attribute of interest to be retrieved
- <table list> : relation names
- <condition> : (Boolean) expression identifying tuples of interest





Relational Database Schema

Simple SQL Queries

| EMPLOYEE | | | | | | | | | |
|----------|-------|-------|------------|-------|---------|-----|--------|----------|-----|
| FNAME | MINIT | LNAME | <u>SSN</u> | BDATE | ADDRESS | SEX | SALARY | SUPERSSN | DNO |

| DEPARTMENT | | | |
|------------|----------------|--------|--------------|
| DNAME | <u>DNUMBER</u> | MGRSSN | MGRSTARTDATE |

| DEPT_LOCATIONS | |
|----------------|------------------|
| <u>DNUMBER</u> | <u>DLOCATION</u> |

| PROJECT | | | |
|---------|----------------|-----------|------|
| PNAME | <u>PNUMBER</u> | PLOCATION | DNUM |

| WORKS_ON | | |
|-------------|------------|-------|
| <u>ESSN</u> | <u>PNO</u> | HOURS |

| DEPENDENT | | | | |
|-------------|-----------------------|-----|-------|--------------|
| <u>ESSN</u> | <u>DEPENDENT_NAME</u> | SEX | BDATE | RELATIONSHIP |

- Retrieve the birthdate and address of the employee whose name is 'John B. Smith'.

```
SELECT  BDATE, ADDRESS
FROM    EMPLOYEE
WHERE   FNAME='John' AND
        MINIT='B' AND LNAME='Smith'
```

- Results may contain duplicate tuples





Joining multiple tables

- Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT FNAME, LNAME, ADDRESS FROM  
EMPLOYEE, DEPARTMENT  
WHERE  
DNAME='Research'  
AND DNUMBER = DNO
```

- For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

```
SELECT  PNUMBER, DNUM, LNAME, BDATE,  
        ADDRESS  
FROM    PROJECT, DEPARTMENT, EMPLOYEE  
WHERE   DNUM=DNUMBER  
        AND MGRSSN=SSN  
        AND PLOCATION='Stafford'
```





Aliases, * and DISTINCT, Empty WHERE-clause

- You can use the same attribute name in different relations
- Prefixing the relation name to the attribute name
- Example: EMPLOYEE.LNAME, DEPARTMENT.DNAME
- ALIASES
 - Sometimes you need to refer same table more than once twice
 - Ex : For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

```
SELECT    E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM      EMPLOYEE E, EMPLOYEE S
WHERE     E.SUPERSSN=S.SSN
```





WHERE-clause

- Missing WHERE-clause = no condition = all tuples
- Ex: Retrieve the SSN values for all employees.

```
SELECT  SSN  
FROM    EMPLOYEE
```

- IMPORTANT: When more than one relation is specified in the FROM-clause this will cause CARTESIAN PRODUCT of rows

```
SELECT  SSN, DNAME  
FROM    EMPLOYEE, DEPARTMENT
```





USE OF *

- * retrieves all the attribute values = all the attributes

Examples:

Q1C: SELECT *
 FROM EMPLOYEE
 WHERE DNO=5

Q1D: SELECT *
 FROM EMPLOYEE, DEPARTMENT
 WHERE DNAME='Research' AND
 DNO=DNUMBER

USE OF DISTINCT

- **DISTINCT** eliminates duplicate
- Example: duplicate SALARY

SELECT SALARY
FROM EMPLOYEE


SELECT **DISTINCT** SALARY
FROM EMPLOYEE



QUERY SET OPERATIONS

- Set operators: (UNION), (MINUS) and (INTERSECT)
- UNION vs UNION ALL (duplicate tuples)
- Queries compatibility :same attributes data type order
- Retrieve all project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

```
(SELECT PNAME
FROM PROJECT, DEPARTMENT, EMPLOYEE
WHERE DNUM=DNUMBER AND MGRSSN=SSN AND LNAME='Smith')
UNION
(SELECT PNAME
FROM PROJECT, WORKS_ON, EMPLOYEE
WHERE PNUMBER=PNO AND ESSN=SSN AND NAME='Smith')
```





NESTING OF QUERIES

- Nested query is a complete select statement (inner) specified within another query (outer query)
- Ex: Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT      FNAME, LNAME, ADDRESS
FROM        EMPLOYEE
WHERE       DNO IN (SELECT DNUMBER
                    FROM DEPARTMENT
                    WHERE DNAME='Research' )
```

- Correlated nested queries:
 - WHERE-clause of a nested query references an attribute in the outer query
 - Ex: Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
SELECT E.FNAME, E.LNAME
FROM      EMPLOYEE AS E
WHERE E.SSN IN (SELECT ESSN
                FROM DEPENDENT
                WHERE ESSN=E.SSN AND
                E.FNAME=DEPENDENT_NAME)
```





THE EXISTS FUNCTION

- EXISTS checks emptiness of nested query is empty (no rows)
- Ex: Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
SELECT    FNAME, LNAME
FROM      EMPLOYEE
WHERE     EXISTS (SELECT *
                  FROM      DEPENDENT
                  WHERE      SSN=ESSN
                  AND FNAME=DEPENDENT_NAME)
```





NULLS IN SQL QUERIES

- SQL uses IS or IS NOT to compare NULLs
- Ex: Retrieve the names of all employees who do not have supervisors.

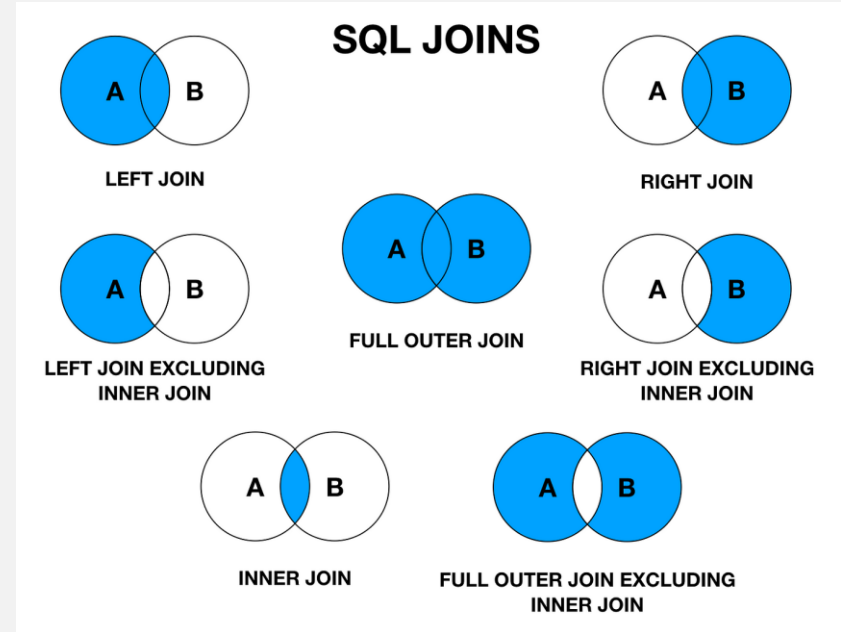
```
SELECT    FNAME, LNAME  
FROM      EMPLOYEE  
WHERE     SUPERSSN IS NULL
```





Joined Relations Feature

- Can specify a "joined relation" in the FROM-clause
- Regular "theta" JOIN, NATURAL JOIN, LEFT OUTER JOIN, RIGHT OUTER JOIN, CROSS JOIN,





Joining Relations

- Examples:

```
SELECT    E.FNAME, E.LNAME,  
          S.FNAME, S.LNAME  
FROM      EMPLOYEE E, EMPLOYEE S  
WHERE     E.SUPERSSN=S.SSN
```

can be written as:

```
SELECT    E.FNAME, E.LNAME,  
          S.FNAME, S.LNAME  
FROM      (EMPLOYEE E LEFT OUTER  
JOIN EMPLOYEES ON  
E.SUPERSSN=S.SSN)
```

- Examples:

```
SELECT FNAME, LNAME, ADDRESS  
FROM EMPLOYEE, DEPARTMENT  
WHERE DNAME='Research' AND  
DNUMBER=DNO
```

could be written as:

```
SELECT FNAME, LNAME, ADDRESS  
FROM (EMPLOYEE JOIN DEPARTMENT  
ON DNUMBER=DNO)  
WHERE DNAME='Research'
```





Joining Relations

- Another Example:

```
SELECT  PNUMBER, DNUM, LNAME, BDATE, ADDRESS  
FROM    (PROJECT JOIN DEPARTMENT ON DNUM=DNUMBER)  
JOIN EMPLOYEE ON MGRSSN=SSN) )  
WHERE   PLOCATION='Stafford'
```





AGGREGATE FUNCTIONS

- Includes COUNT, SUM, MAX, MIN, and AVG
- Query: Find the maximum salary, the minimum salary, and the average salary among all employees.

```
SELECT  MAX(SALARY), MIN(SALARY), AVG(SALARY)
FROM    EMPLOYEE
```

- Query: Find the maximum salary, the minimum salary, and the average salary among employees who work for the 'Research' department.

```
SELECT  MAX(SALARY), MIN(SALARY), AVG(SALARY)
FROM    EMPLOYEE, DEPARTMENT
WHERE   DNO=DNUMBER AND DNAME='Research'
```





AGGREGATE FUNCTIONS (contd.)

- Queries: Retrieve the total number of employees in the company and the number of employees in the 'Research' department .

```
SELECT    COUNT (*)  
FROM      EMPLOYEE
```

```
SELECT    COUNT (*)  
FROM      EMPLOYEE, DEPARTMENT  
WHERE     DNO=DNUMBER AND DNAME='Research'
```





GROUPING

- Applies aggregate functions to subgroups of rows having the same value for the grouping attribute(s)
- Aggregate function is applied to each subgroup independently
- GROUP BY-clause lists grouping attributes (must appear in the SELECT-clause)
- Query: For each department, retrieve the department number, the number of employees in the department, and their average salary.

```
SELECT    DNO, COUNT (*), AVG (SALARY)
FROM      EMPLOYEE
GROUP BY DNO
```

- Query : For each project, retrieve the project number, project name, and the number of employees who work on that project.

```
SELECT    PNUMBER, PNAME, COUNT (*)
FROM      PROJECT, WORKS_ON
WHERE     PNUMBER=PNO
GROUP BY PNUMBER, PNAME
```





THE HAVING-CLAUSE

- HAVING-clause retrieves functions for only groups satisfying certain conditions
- Query : For each project on which more than two employees work, retrieve the project number, project name, and the number of employees who work on that project.

```
SELECT          PNUMBER, PNAME, COUNT(*)  
FROM            PROJECT, WORKS_ON  
WHERE           PNUMBER=PNO  
GROUP BY        PNUMBER, PNAME  
HAVING          COUNT (*) > 2
```





SUBSTRING COMPARISON

- LIKE compares partial strings
 - '%' (or '*' in some implementations) replaces an arbitrary number of characters,
 - '_' replaces a single arbitrary character
- Query: Retrieve all employees whose address is in Houston, Texas. Here, the value of the ADDRESS attribute must contain the substring 'Houston,TX' in it.

```
SELECT FNAME, LNAME  
FROM   EMPLOYEE  
WHERE  ADDRESS LIKE '%Houston,TX%'
```





ARITHMETIC OPERATIONS

- Arithmetic operators '+', '-', '*', and '/'
- Query : Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

```
SELECT FNAME, LNAME, 1.1*SALARY  
FROM   EMPLOYEE, WORKS_ON, PROJECT  
WHERE  SSN=ESSN AND PNO=PNUMBER AND PNAME='ProductX'
```





ORDER BY

- ORDER BY clause sorts query result based on some attributes values
- Specify the keyword DESC or ASC (default)
- Query : Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name.

```
SELECT      DNAME, LNAME, FNAME, PNAME
FROM        DEPARTMENT, EMPLOYEE, WORKS_ON, PROJECT
WHERE       DNUMBER=DNO AND SSN=ESSN AND PNO=PNUMBER
ORDER BY    DNAME, LNAME
```





Summary of SQL Queries

| | |
|-----------|--------------------------|
| SELECT | <attribute list> |
| FROM | <table list> |
| [WHERE | <condition>] |
| [GROUP BY | <grouping attribute(s)>] |
| [HAVING | <group condition>] |
| [ORDER BY | <attribute list>] |





DML

- Three SQL commands : INSERT, DELETE, and UPDATE
- INSERT: add one or more rows to a table, values should be listed in the same order as the attributes were specified in the CREATE TABLE command

Example: **INSERT INTO EMPLOYEE**

VALUES ('Richard','K','Marini', '653298653', '30-DEC-52','98 Oak Forest,Katy,TX', 'M', 37000,'987654321', 4)

- Alternative: specify explicitly the attribute names

Example: **INSERT INTO EMPLOYEE (FNAME, LNAME, SSN)**

VALUES ('Richard', 'Marini', '653298653')





DELETE

- Removes rows from a table
 - Can have WHERE-clause to select the tuples to be deleted
 - Omitting WHERE-clause specifies that all rows tuples are to be deleted
- Examples:

```
DELETE FROM EMPLOYEE WHERE LNAME='Brown'
```

```
DELETE FROM EMPLOYEE WHERE SSN='123456789'
```

```
DELETE FROM EMPLOYEE WHERE DNO IN  
(SELECT   DNUMBER  
  FROM    DEPARTMENT WHERE DNAME='Research')
```

```
DELETE FROM EMPLOYEE
```





UPDATE

- Modify values of one or more selected rows satisfying WHERE-clause if exists while enforcing referential integrity
- Example: Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively.

UPDATE PROJECT

SET PLOCATION = 'Bellaire', DNUM = 5
WHERE PNUMBER=10

- Example: Give all employees in the 'Research' department a 10% raise in salary.

UPDATE EMPLOYEE
SET SALARY = SALARY *1.1
WHERE DNO IN (SELECT DNUMBER
FROM DEPARTMENT
WHERE DNAME='Research')





Normalization of Relations

- **Normalization:** The process of decomposing unsatisfactory "bad" relations by breaking up their attributes into smaller relations so the resulting designs are of high quality and meet the desirable properties
- **Normal form:** Condition using keys and FDs of a relation to certify whether a relation schema is in a particular normal form
- 2NF, 3NF, BCNF , 4NF
 - based on keys and FDs of a relation schema
- **Denormalization:** The process of storing the join of higher normal form relations as a base relation—which is in a lower normal form





Definitions of Keys and Attributes Participating in Keys

- A **superkey** of a relation schema $R = \{A_1, A_2, \dots, A_n\}$ is a set of attributes S subset-of R with the property that no two tuples t_1 and t_2 in any legal relation state r of R will have $t_1[S] = t_2[S]$
- A **key** K is a **superkey** with the additional property that removal of any attribute from K will cause K not to be a superkey any more.
- If a relation schema has more than one key, each is called a **candidate key**.
 - One of the candidate keys is arbitrarily designated to be the **primary key**, and the others are called **secondary keys**.
- A **Prime attribute** must be a member of some candidate key
- A **Nonprime attribute** is not a prime attribute—that is, it is not a member of any candidate key.

First Normal Form

- Disallows
 - Composite attributes
 - Multivalued attributes
 - Nested relations;** attributes whose values for an individual tuple are non-atomic

(a)

| DEPARTMENT | | | |
|------------|---------|----------|------------|
| Dname | Dnumber | Dmgr_ssn | Dlocations |
| | | | |

(b)

| DEPARTMENT | | | |
|----------------|---------|-----------|--------------------------------|
| Dname | Dnumber | Dmgr_ssn | Dlocations |
| Research | 5 | 333445555 | {Bellaire, Sugarland, Houston} |
| Administration | 4 | 987654321 | {Stafford} |
| Headquarters | 1 | 888665555 | {Houston} |

(c)

| DEPARTMENT | | | |
|----------------|---------|-----------|-----------|
| Dname | Dnumber | Dmgr_ssn | Dlocation |
| Research | 5 | 333445555 | Bellaire |
| Research | 5 | 333445555 | Sugarland |
| Research | 5 | 333445555 | Houston |
| Administration | 4 | 987654321 | Stafford |
| Headquarters | 1 | 888665555 | Houston |

a)

| EMP_PROJ | | Projs | |
|----------|-------|---------|-------|
| Ssn | Ename | Pnumber | Hours |
| | | | |

b)

| EMP_PROJ | | | |
|-----------|----------------------|---------|-------|
| Ssn | Ename | Pnumber | Hours |
| 123456789 | Smith, John B. | 1 | 32.5 |
| | | 2 | 7.5 |
| 666884444 | Narayan, Ramesh K. | 3 | 40.0 |
| 453453453 | English, Joyce A. | 1 | 20.0 |
| | | 2 | 20.0 |
| 333445555 | Wong, Franklin T. | 2 | 10.0 |
| | | 3 | 10.0 |
| | | 10 | 10.0 |
| | | 20 | 10.0 |
| 999887777 | Zelaya, Alicia J. | 30 | 30.0 |
| | | 10 | 10.0 |
| 987987987 | Jabbar, Ahmad V. | 10 | 35.0 |
| | | 30 | 5.0 |
| 987654321 | Wallace, Jennifer S. | 30 | 20.0 |
| | | 20 | 15.0 |
| 888665555 | Borg, James E. | 20 | NULL |

c)

| EMP_PROJ1 | |
|-----------|-------|
| Ssn | Ename |
| | |

| EMP_PROJ2 | | |
|-----------|---------|-------|
| Ssn | Pnumber | Hours |
| | | |





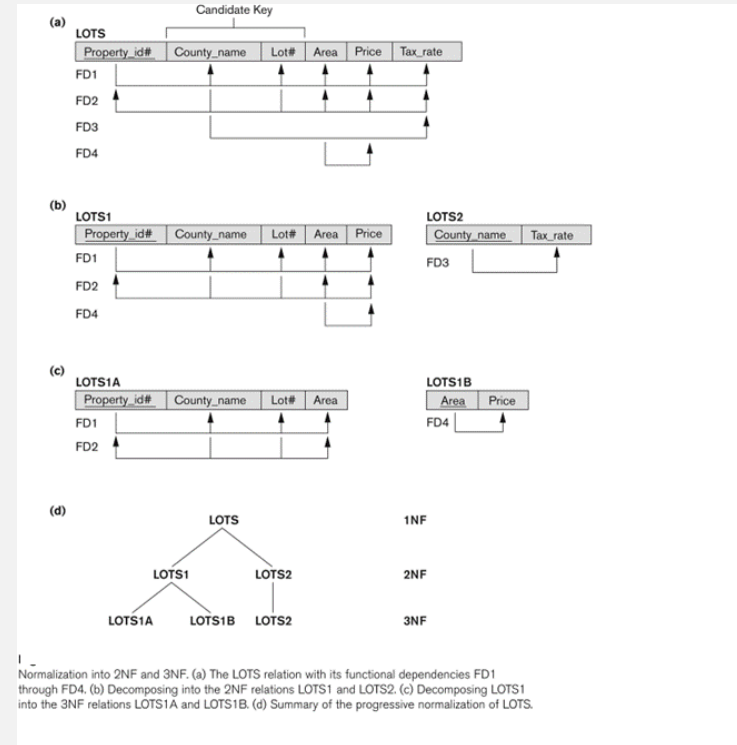
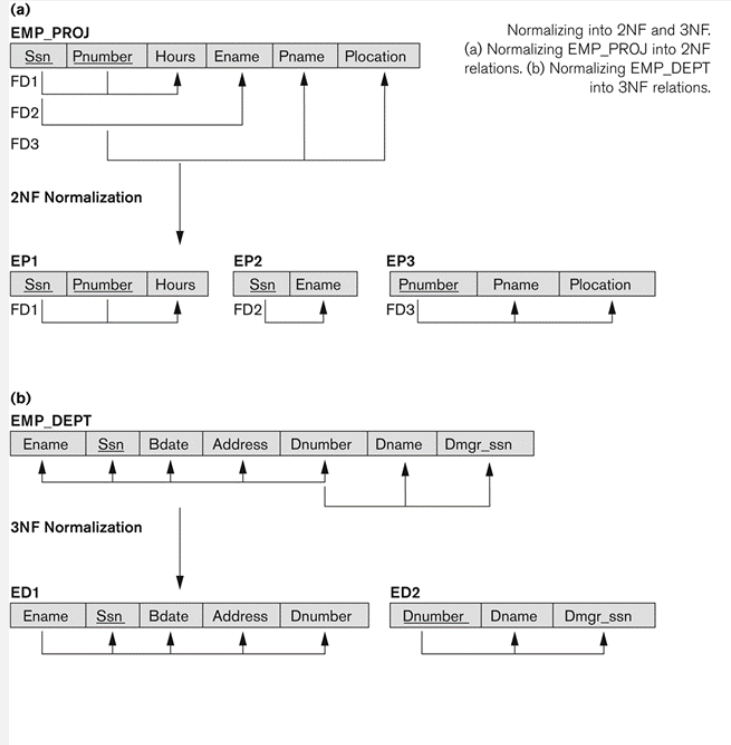
Second Normal Form

- Uses the concepts of FDs, primary key
- Definitions
 - **Prime attribute**: An attribute that is member of the primary key K
 - **Full functional dependency**: a FD $Y \rightarrow Z$ where removal of any attribute from Y means the FD does not hold any more
- Examples:
 - $\{SSN, PNUMBER\} \rightarrow HOURS$ is a full FD since neither $SSN \rightarrow HOURS$ nor $PNUMBER \rightarrow HOURS$ hold
 - $\{SSN, PNUMBER\} \rightarrow ENAME$ is not a full FD (it is called a partial dependency) since $SSN \rightarrow ENAME$ also holds
- A relation schema R is in **second normal form (2NF)** if every non-prime attribute A in R is fully functionally dependent on the primary key
- R can be decomposed into 2NF relations via the process of 2NF normalization





Normalizing into 2NF and 3NF



Third Normal Form

- Definition:
 - **Transitive functional dependency:** a FD $X \rightarrow Z$ that can be derived from two FDs $X \rightarrow Y$ and $Y \rightarrow Z$
- Examples:
 - $SSN \rightarrow DMGRSSN$ is a transitive FD
 - Since $SSN \rightarrow DNUMBER$ and $DNUMBER \rightarrow DMGRSSN$ hold
 - $SSN \rightarrow ENAME$ is non-transitive
 - Since there is no set of attributes X where $SSN \rightarrow X$ and $X \rightarrow ENAME$
- A relation schema R is in third normal form (3NF) if it is in 2NF and no non-prime attribute A in R is transitively dependent on the primary key
- R can be decomposed into 3NF relations via the process of 3NF normalization
- NOTE:
 - In $X \rightarrow Y$ and $Y \rightarrow Z$, with X as the primary key, we consider this a problem only if Y is not a candidate key.
 - When Y is a candidate key, there is no problem with the transitive dependency .
 - E.g., Consider $EMP(SSN, Emp\#, Salary)$.
 - Here, $SSN \rightarrow Emp\# \rightarrow Salary$ and $Emp\#$ is a candidate key.





Normal Forms Defined Informally

1st normal form

- All attributes depend on the **key**

2nd normal form

- All attributes depend on the **whole key**

3rd normal form

- All attributes depend on **nothing but the key**





Data Warehouses

W. H Inmon characterized a data warehouse as:

- **“A subject-oriented, integrated, nonvolatile, time-variant collection of data in support of management’s decisions.”**

Data warehouses have the distinguishing characteristic that they are mainly intended for decision support applications.

- Traditional databases are transactional.

Applications that data warehouse supports are:

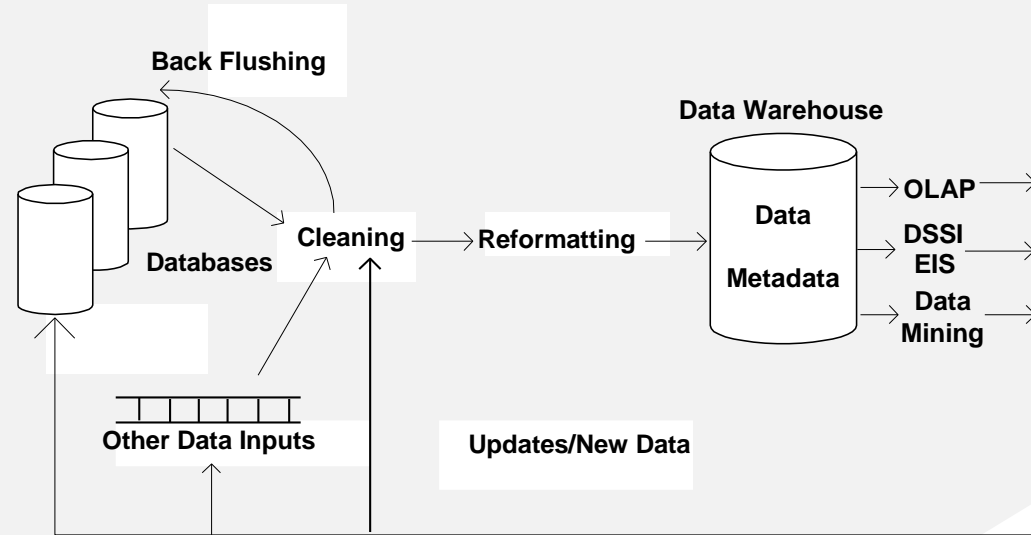
- **OLAP** (Online Analytical Processing) is a term used to describe the analysis of complex data from the data warehouse.
- **DSS** (Decision Support Systems) also known as EIS (Executive Information Systems) supports organization’s leading decision makers for making complex and important decisions.
- **Data Mining** is used for knowledge discovery, the process of searching data for unanticipated new knowledge.





Conceptual Structure of Data Warehouse

- Data Warehouse processing involves
- Cleaning and reformatting of data
 - OLAP
 - Data Mining





Comparison with Traditional Databases

Data Warehouses are mainly optimized for appropriate data access.

- **Traditional databases are transactional and are optimized for both access mechanisms and integrity assurance measures.**

Data warehouses emphasize more on historical data as their main purpose is to support time-series and trend analysis.

Compared with transactional databases, data warehouses are nonvolatile.

In transactional databases transaction is the mechanism change to the database. By contrast information in data warehouse is relatively coarse grained and refresh policy is carefully chosen, usually incremental.





Characteristics of Data Warehouses

- Multidimensional conceptual view
- Generic dimensionality
- Unlimited dimensions and aggregation levels
- Unrestricted cross-dimensional operations
- Dynamic sparse matrix handling
- Client-server architecture
- Multi-user support
- Accessibility
- Transparency
- Intuitive data manipulation
- Consistent reporting performance
- Flexible reporting





Classification of Data Warehouses

Generally, Data Warehouses are an order of magnitude larger than the source databases. The sheer volume of data is an issue, based on which Data Warehouses could be classified as follows.

- **Enterprise-wide data warehouses**
 - They are huge projects requiring massive investment of time and resources.
- **Virtual data warehouses**
 - They provide views of operational databases that are materialized for efficient access.
- **Data marts**
 - These are generally targeted to a subset of organization, such as a department, and are more tightly focused.





Data Modeling for Data Warehouses

Traditional Databases generally deal with two-dimensional data (similar to a spread sheet).

- However, querying performance in a multi-dimensional data storage model is much more efficient.

Data warehouses can take advantage of this feature as generally these are

- Non volatile
- The degree of predictability of the analysis that will be performed on them is high.






Data Modeling for Data Warehouses

Example of Two- Dimensional vs. Multi- Dimensional

Two Dimensional Model

| | | REGION | | |
|---------------------------------|------|--------|------|------|
| | | REG1 | REG2 | REG3 |
| P R O D U C T | P123 | | | |
| | P124 | | | |
| | P125 | | | |
| | P126 | | | |
| | ⋮ | | | |





Data Modeling for Data Warehouses

Advantages of a multi-dimensional model

- Multi-dimensional models lend themselves readily to hierarchical views in what is known as roll-up display and drill-down display.
- The data can be directly queried in any combination of dimensions, bypassing complex database queries.

Multi-dimensional schemas are specified using:

- **Dimension table**
 - It consists of tuples of attributes of the dimension.
- **Fact table**
 - Each tuple is a recorded fact. This fact contains some measured or observed variable (s) and identifies it with pointers to dimension tables. The fact table contains the data, and the dimensions to identify each tuple in the data.

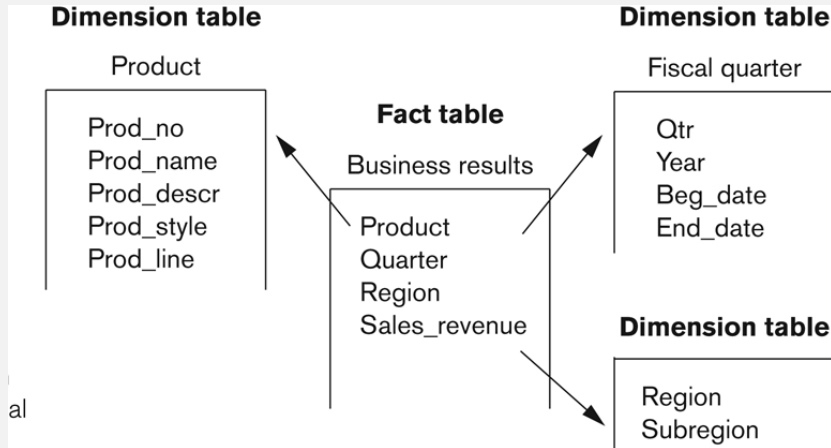




Multi-dimensional Schemas

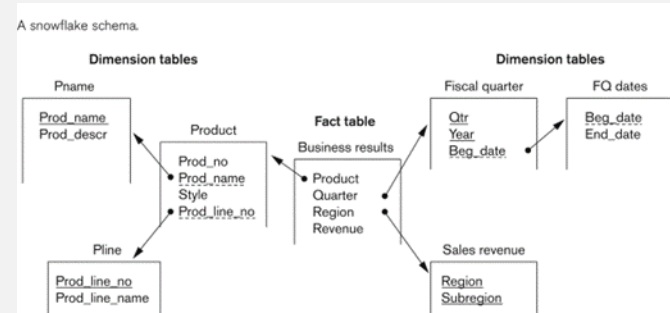
Star schema:

- Consists of a fact table with a single table for each dimension



Snowflake Schema:

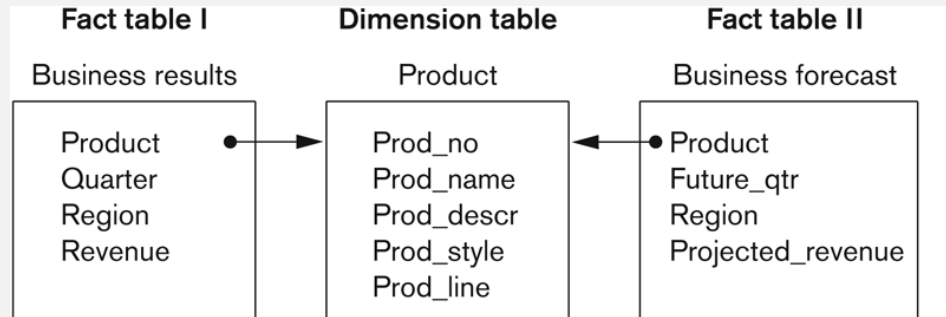
- It is a variation of star schema, in which the dimensional tables from a star schema are organized into a hierarchy by normalizing them.





Multi-dimensional Schemas

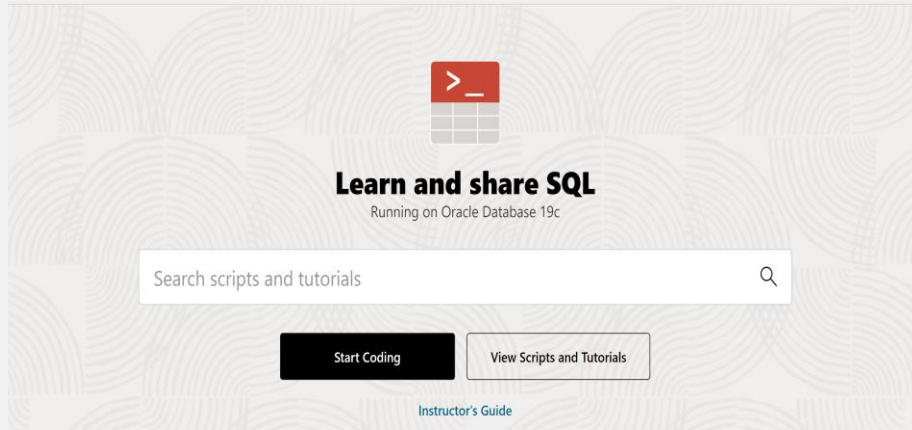
- **Fact Constellation**
 - Fact constellation is a set of tables that share some dimension tables. However, fact constellations limit the possible queries for the warehouse.





Oracle Live SQL

<https://livesql.oracle.com/>



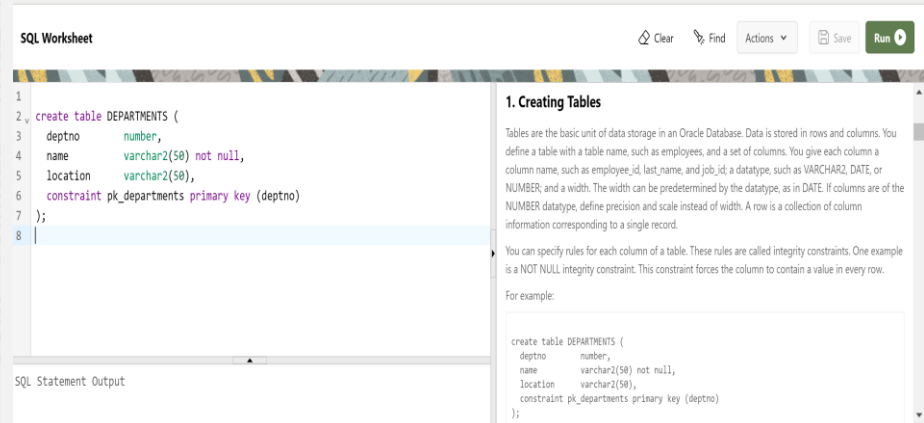
The landing page features a red terminal icon with a white prompt character. Below it, the text "Learn and share SQL" is prominently displayed, followed by "Running on Oracle Database 19c". A search bar with the placeholder "Search scripts and tutorials" and a magnifying glass icon is positioned below the text. At the bottom, there are two buttons: "Start Coding" and "View Scripts and Tutorials". A link for "Instructor's Guide" is located at the bottom right.

Learn and share SQL
Running on Oracle Database 19c

Search scripts and tutorials

[Start Coding](#) [View Scripts and Tutorials](#)

[Instructor's Guide](#)



The interface shows an "SQL Worksheet" with a toolbar at the top containing "Clear", "Find", "Actions", "Save", and "Run" buttons. The main area contains a SQL script for creating a table named DEPARTMENTS. The script is as follows:

```
1  
2 create table DEPARTMENTS (  
3   deptno    number,  
4   name      varchar2(50) not null,  
5   location  varchar2(50),  
6   constraint pk_departments primary key (deptno)  
7 );  
8
```

On the right side, there is a section titled "1. Creating Tables" with explanatory text about tables and integrity constraints. Below this text, the same SQL script is shown in a code block.

1. Creating Tables

Tables are the basic unit of data storage in an Oracle Database. Data is stored in rows and columns. You define a table with a table name, such as employees, and a set of columns. You give each column a column name, such as employee_id, last_name, and job_id; a datatype, such as VARCHAR2, DATE, or NUMBER; and a width. The width can be predetermined by the datatype, as in DATE. If columns are of the NUMBER datatype, define precision and scale instead of width. A row is a collection of column information corresponding to a single record.

You can specify rules for each column of a table. These rules are called integrity constraints. One example is a NOT NULL integrity constraint. This constraint forces the column to contain a value in every row.

For example:

```
create table DEPARTMENTS (  
  deptno    number,  
  name      varchar2(50) not null,  
  location  varchar2(50),  
  constraint pk_departments primary key (deptno)  
);
```

SQL Statement Output





Oracle Live SQL

Tutorial

Introduction to SQL

This tutorial provides an introduction to the Structured Query Language (SQL), learn how to create t...

create table, create, select, insert, update, delete, drop, drop table, recycle bin, purge

Tutorial

Sorting and Limiting Rows: Databases for Developers

An introduction to sorting data with order by and restricting rows to the top N.

order by, fetch first

Tutorial

Updating table data

This tutorial demonstrates different variations of the UPDATE statement. It includes examples of bas...

Tutorial

Aggregating Rows: Databases for Developers

An introduction to how to summarise data using aggregate functions and group by.

group by, rollup, cube, count, sum

Tutorial

Joining Tables: Databases for Developers

An introduction to the join types available in Oracle Database.

join, inner join, outer join, cross join

Tutorial

Subqueries: Databases for Developers

An introduction to using subqueries in Oracle Database

with clause, in, exists

Tutorial

Union, Minus, and Intersect: Databases for Developers

An introduction to the set operators, union, minus, and intersect

union, minus, intersect

Tutorial

Analytic Functions: Databases for Developers

An introduction to analytic functions.

analytics



Thank You



SDAIA
الهيئة السعودية للبيانات
والذكاء الاصطناعي
Saudi Data & AI Authority